

Practical 1

AIM :- Implement Linear Regression (Diabetes Dataset)

THEORY :-

Logistic regression is a classification model in machine learning, extensively used in clinical analysis. It uses probabilistic estimations which helps in understanding the relationship between the dependent variable and one or more independent variables. Diabetes, being one of the most common diseases around the world, when detected early, may prevent the progression of the disease and avoid other complications. In this work, we design a prediction model, that predicts whether a patient has diabetes, based on certain diagnostic measurements included in the dataset, and explore various techniques to boost the performance and accuracy.

CODE / OUTPUT :-

```
# Import Dependencies

import numpy as np

import matplotlib.pyplot as plt

from sklearn import datasets,linear_model,metrics
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

import seaborn as sns

# Load the diabetes dataset

diabetes=datasets.load_diabetes()

# X - feature vectors

# y - Target values

X=diabetes.data

y=diabetes.target

# splitting X and y into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4,
```

```
random_state=1)
```

```
# Create linear regression object
```

```
lin_reg=linear_model.LinearRegression()
```

```
# Train the model using train and test data
```

```
lin_reg.fit(X_train,y_train)
```

```
Out[6]: LinearRegression()
```

```
# Predict values for X_test data
```

```
predicted = lin_reg.predict(X_test)
```

```
# Regression coefficients
```

```
print('\n Coefficients are:\n',lin_reg.coef_)
```

```
# Intercept
```

```
print('\nIntercept : ',lin_reg.intercept_)
```

```
# variance score: 1 means perfect prediction
```

```
print('Variance score: ',lin_reg.score(X_test, y_test))
```

```
Coefficients are:
```

```
[ -59.73663337 -215.62170919  599.92621335  291.96724002 -829.65206295  
 544.63994617  164.85191153  224.2392528   768.94426062   70.84982207]
```

```
Intercept : 152.89009028286725
```

```
Variance score: 0.4160439011127657
```

```
# Mean Squared Error
```

```
print("Mean squared error: %.2f\n"
```

```
      % mean_squared_error(y_test, predicted))
```

```
# Original data of X_test
```

```
expected = y_test
```

```
Mean squared error: 2962.93
```

Plot a graph for expected and predicted values

```
plt.title('Linear Regression ( DIABETES Dataset)')  
plt.scatter(expected,predicted,c='b',marker='.',s=36)  
plt.plot(np.linspace(0, 330, 100),np.linspace(0, 330, 100), '--r', linewidth=2)  
plt.show()
```

